

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
import plotly.express as px
import plotly.io as pio

df=pd.read_excel("/content/Churn_Modelling.xlsx")

```

df

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenur
0	1.0	15634602.0	Hargrave	619.0	France	Female	42.0	2.
1	2.0	15647311.0	Hill	608.0	Spain	Female	41.0	1.
2	3.0	15619304.0	Onio	502.0	France	Female	42.0	8.
3	4.0	15701354.0	Boni	699.0	France	Female	39.0	1.
4	5.0	15737888.0	Mitchell	850.0	Spain	Female	43.0	2.
...	...	...	...	...	...	...	...	...
9995	9996.0	15606229.0	Obijiaku	771.0	France	Male	39.0	5.
9996	9997.0	15569892.0	Johnstone	516.0	France	Male	35.0	10.
9997	9998.0	15584532.0	Liu	709.0	France	Female	36.0	7.
9998	9999.0	15682355.0	Sabbatini	772.0	Germany	Male	42.0	3.
9999	10000.0	15628319.0	Walker	792.0	France	Female	28.0	4.

10000 rows × 14 columns

df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   RowNumber        10000 non-null   float64
 1   CustomerId       10000 non-null   float64
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   float64
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   float64
 7   Tenure            10000 non-null   float64
 8   Balance           10000 non-null   float64
 9   NumOfProducts     10000 non-null   float64
 10  HasCrCard        10000 non-null   float64
 11  IsActiveMember    10000 non-null   float64
 12  EstimatedSalary   10000 non-null   float64
 13  Exited           10000 non-null   float64
dtypes: float64(11), object(3)
memory usage: 1.1+ MB

```

df.isnull().sum()

```

RowNumber      0
CustomerId     0
Surname        0
CreditScore    0
Geography      0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited          0
dtype: int64

```

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Ba
<b>count</b>	10000.00000	1.000000e+04	10000.00000	10000.00000	10000.00000	10000.0
<b>mean</b>	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.8
<b>std</b>	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.4
<b>min</b>	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.0
<b>25%</b>	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.0
<b>50%</b>	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.5
<b>75%</b>	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.2
<b>max</b>	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.0

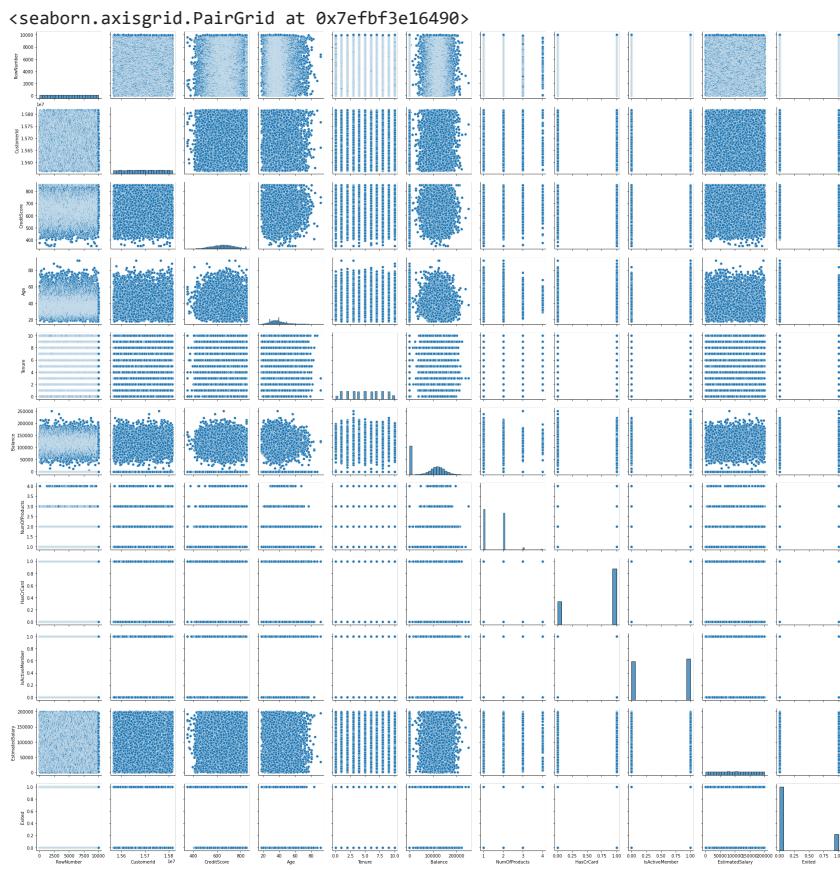
```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   RowNumber        10000 non-null   float64
 1   CustomerId       10000 non-null   float64
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   float64
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   float64
 7   Tenure            10000 non-null   float64
 8   Balance           10000 non-null   float64
 9   NumOfProducts     10000 non-null   float64
 10  HasCrCard        10000 non-null   float64
 11  IsActiveMember   10000 non-null   float64
 12  EstimatedSalary  10000 non-null   float64
 13  Exited            10000 non-null   float64
dtypes: float64(11), object(3)
memory usage: 1.1+ MB
```

```
df.corr()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
<b>RowNumber</b>	1.00000	0.004202	0.005840	0.000783	-0.006495	-0.009067
<b>CustomerId</b>	0.004202	1.00000	0.005308	0.009497	-0.014883	-0.012419
<b>CreditScore</b>	0.005840	0.005308	1.00000	-0.003965	0.000842	0.006268
<b>Age</b>	0.000783	0.009497	-0.003965	1.00000	-0.009997	0.028308
<b>Tenure</b>	-0.006495	-0.014883	0.000842	-0.009997	1.00000	-0.012254
<b>Balance</b>	-0.009067	-0.012419	0.006268	0.028308	-0.012254	1.00000
<b>NumOfProducts</b>	0.007246	0.016972	0.012238	-0.030680	0.013444	-0.304180
<b>HasCrCard</b>	0.000599	-0.014025	-0.005458	-0.011721	0.022583	-0.014858
<b>IsActiveMember</b>	0.012044	0.001665	0.025651	0.085472	-0.028362	-0.010084
<b>EstimatedSalary</b>	-0.005988	0.015271	-0.001384	-0.007201	0.007784	0.012797
<b>Exited</b>	-0.016571	-0.006248	-0.027094	0.285323	-0.014001	0.118533

```
sns.pairplot(df)
```



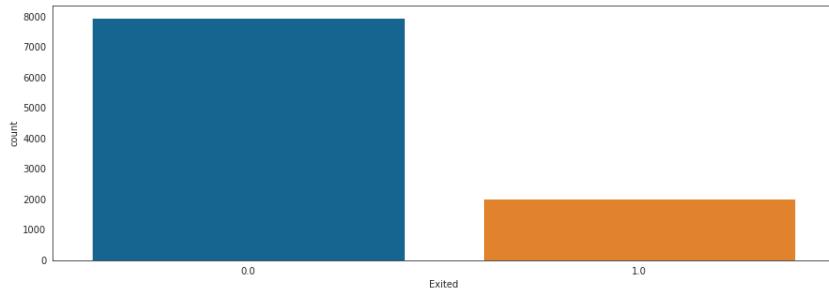
```
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams["figure.figsize"] = (15,5)
sns.set_style('white')
```

```
plt.style.use('tableau-colorblind10')
#plt.xkcd()
```

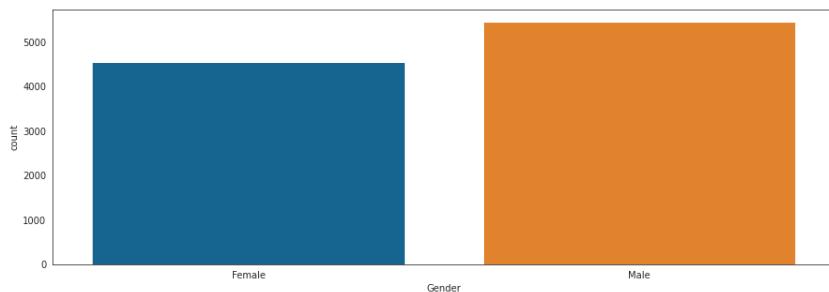
```
plt.style.available
```

```
['Solarize_Light2',
 '_classic_test_patch',
 'bmh',
 'classic',
 'dark_background',
 'fast',
 'fivethirtyeight',
 'ggplot',
 'grayscale',
 'seaborn',
 'seaborn-bright',
 'seaborn-colorblind',
 'seaborn-dark',
 'seaborn-dark-palette',
 'seaborn-darkgrid',
 'seaborn-deep',
 'seaborn-muted',
 'seaborn-notebook',
 'seaborn-paper',
 'seaborn-pastel',
 'seaborn-poster',
 'seaborn-talk',
 'seaborn-ticks',
 'seaborn-white',
 'seaborn-whitegrid',
 'tableau-colorblind10']
```

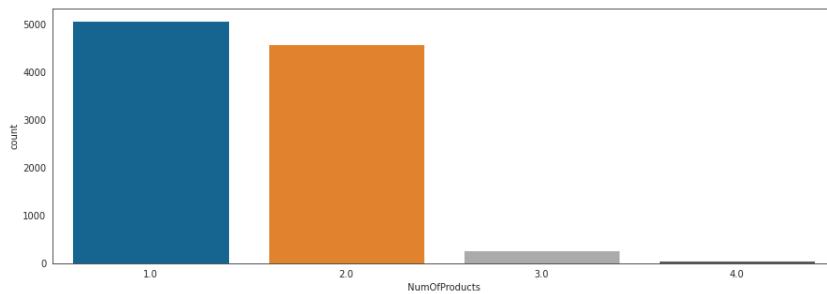
```
ax = sns.countplot('Exited', data=df) #Imbalanced Dataset
```



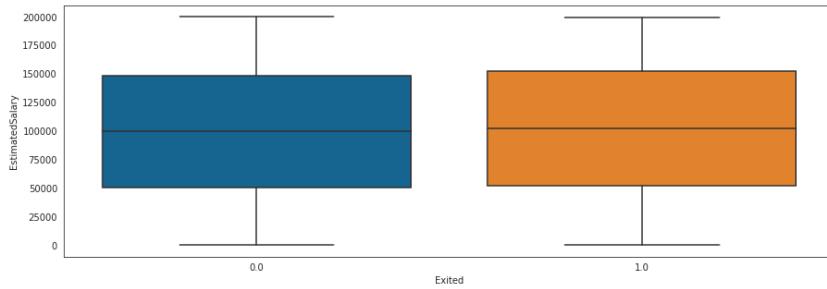
```
ax = sns.countplot('Gender', data=df)
```



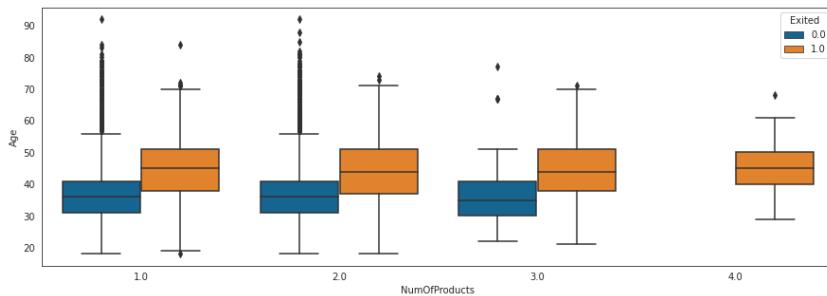
```
ax = sns.countplot('NumOfProducts', data=df)
```



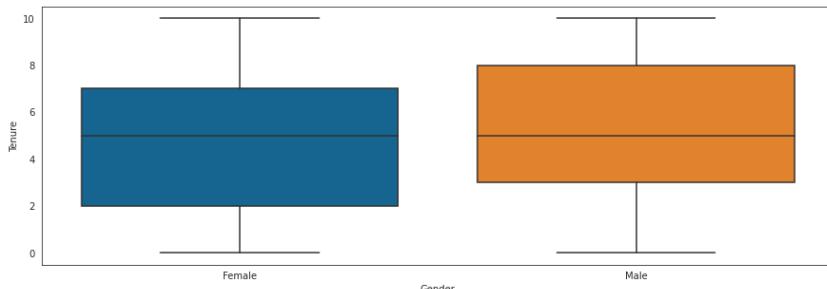
```
ax = sns.boxplot(x='Exited',y='EstimatedSalary',data=df)
```



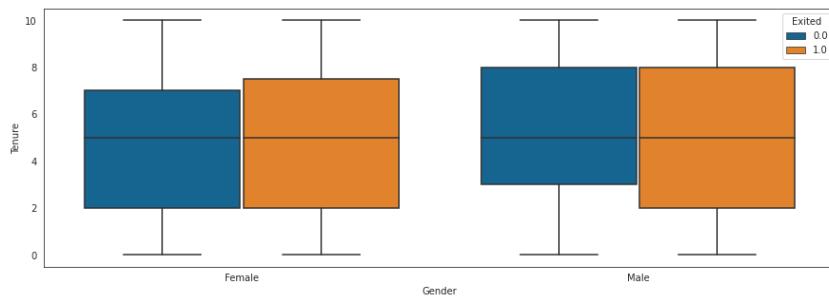
```
ax = sns.boxplot(x='NumOfProducts',y='Age',hue='Exited',data=df)
```



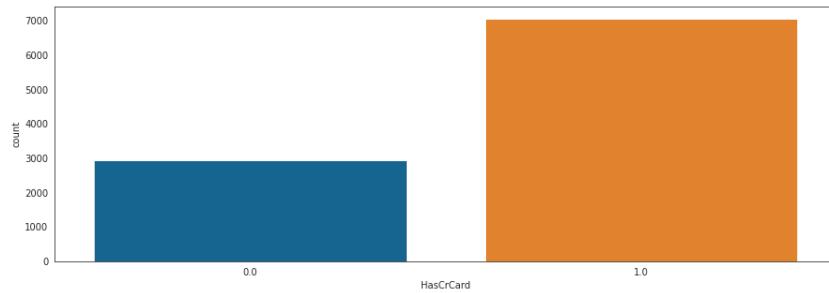
```
ax = sns.boxplot(x='Gender',y='Tenure',data=df)
```



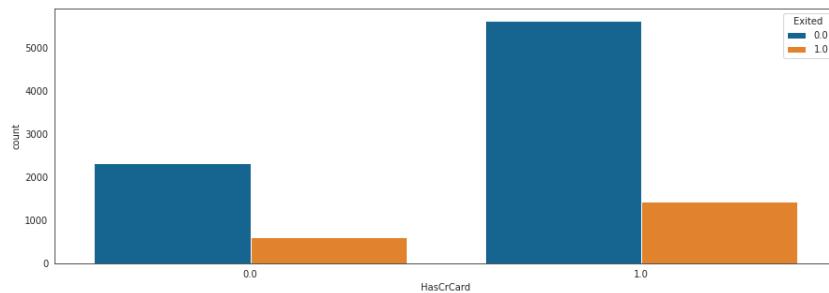
```
ax = sns.boxplot(x='Gender',y='Tenure',hue='Exited',data=df)
```



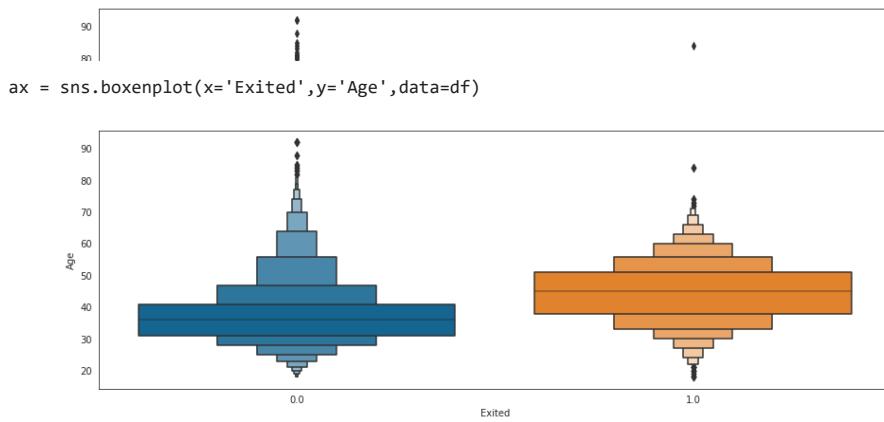
```
#HasCrCard
ax = sns.countplot('HasCrCard', data=df)
```



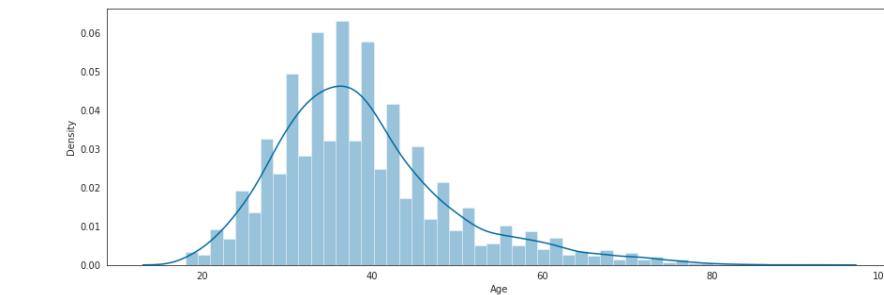
```
ax = sns.countplot('HasCrCard', hue='Exited', data=df)
```



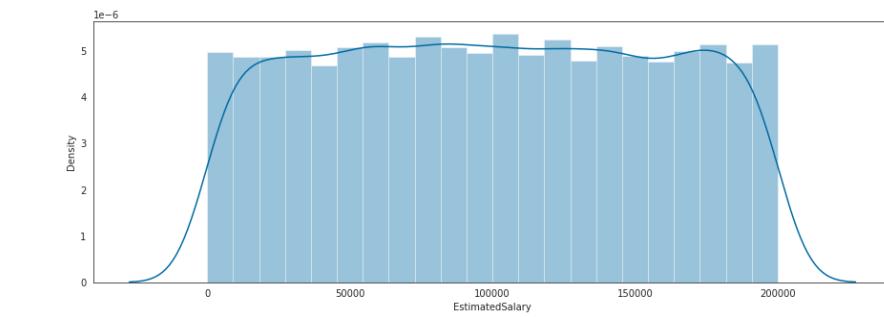
```
ax = sns.boxplot(x='Exited', y='Age', data=df)
```



```
ax = sns.boxenplot(x='Exited',y='Age',data=df)
```



```
#EstimatedSalary
ax = sns.distplot(df.EstimatedSalary)
```

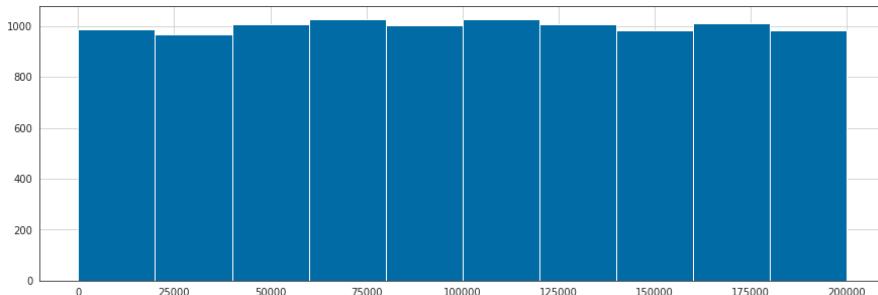


```
df.EstimatedSalary.describe()
```

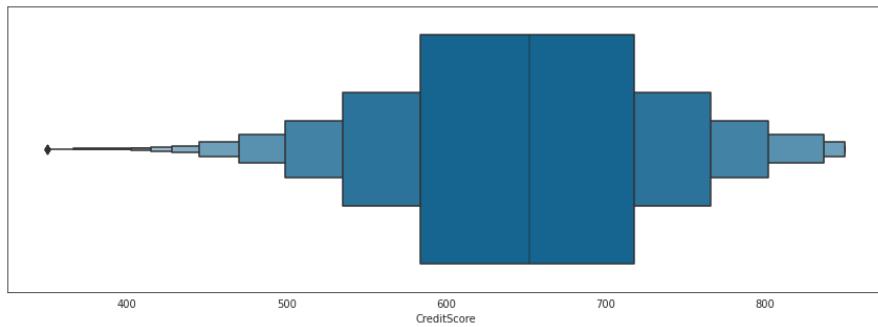
count	10000.000000
mean	100090.239881
std	57510.492818
min	11.580000
25%	51002.110000
50%	100193.915000
75%	149388.247500
max	199992.480000

Name: EstimatedSalary, dtype: float64

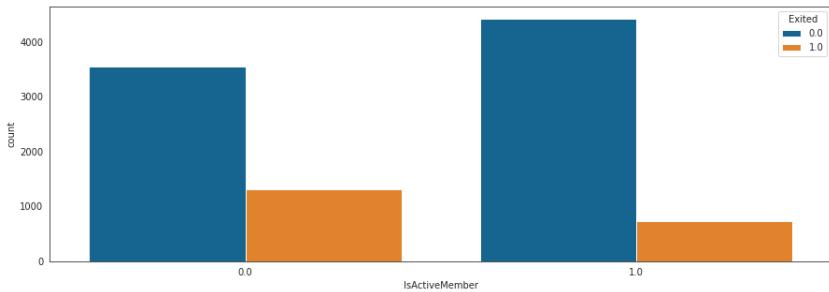
```
ax = df['EstimatedSalary'].hist()
```



```
ax = sns.boxenplot('CreditScore', data=df)
```

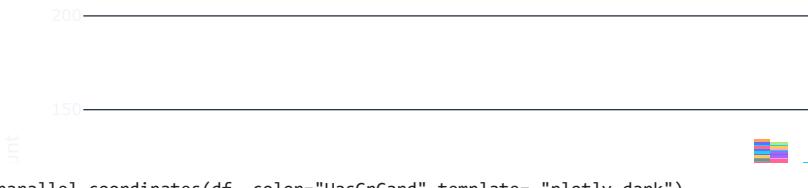


```
ax = sns.countplot('IsActiveMember', hue='Exited', data=df)
```

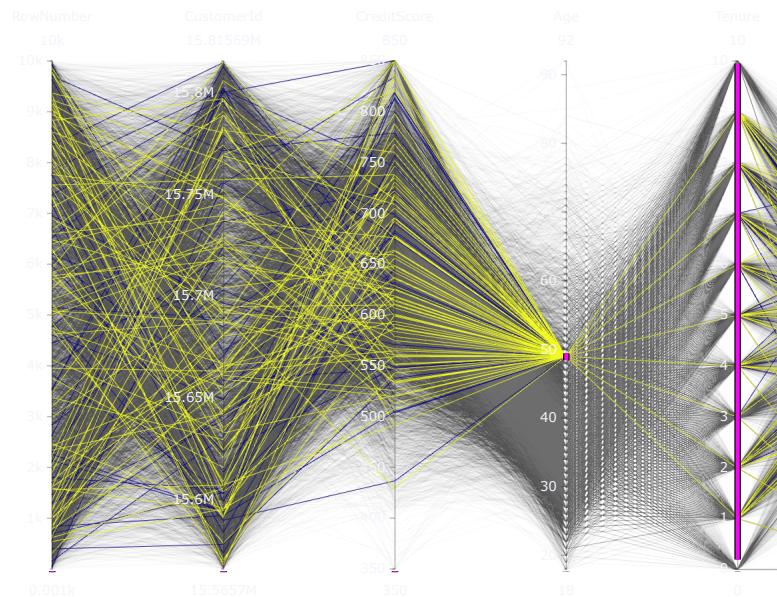


```
ax= px.histogram(df,x= "CreditScore", template= "plotly_dark",color= "Balance",title='credit distribution')
ax.show()
```

## credit distribution



```
ax= px.parallel_coordinates(df, color="HasCrCard",template= "plotly_dark")
ax.show()
```



```
df.drop(["Gender"],axis=1,inplace=True)
```

```
#splittig x and y for predictioions
```

```
x=df.iloc[:,1:-1]
```

```
x
```

```

CustomerID CreditScore Age Tenure Balance NumOfProducts HasCrCard I
0 15634602.0 619.0 42.0 2.0 0.00 1.0 1.0
2 15619304.0 502.0 42.0 8.0 159660.80 3.0 1.0
y
0 1.0
1 0.0
2 1.0
3 0.0
4 0.0
...
9995 0.0
9996 0.0
9997 1.0
9998 1.0
9999 0.0
Name: Exited, Length: 10000, dtype: float64
10000 rows x 9 columns
#splitting the model in x,y train and test

from sklearn.model_selection import train_test_split

xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.30,random_state=1)

from sklearn.linear_model import LogisticRegression

logreg=LogisticRegression()
logreg.fit(xtrain,ytrain)
ypred=logreg.predict(xtest)

from sklearn.metrics import classification_report

print(classification_report(ytest,ypred))

precision recall f1-score support
0.0 0.79 1.00 0.88 2373
1.0 0.00 0.00 0.00 627
accuracy 0.79 3000
macro avg 0.40 0.50 0.44 3000
weighted avg 0.63 0.79 0.70 3000

#feature scaling

from sklearn.preprocessing import StandardScaler

sc=StandardScaler()
xtrain=sc.fit_transform(xtrain)

xtest=sc.transform(xtest)

#hypertuning using solver parameter

logreg=LogisticRegression(solver="saga")
logreg.fit(xtrain,ytrain)
ypred=logreg.predict(xtest)

print(classification_report(ytest,ypred))

precision recall f1-score support
0.0 0.82 0.98 0.89 2373
1.0 0.65 0.17 0.27 627

```

accuracy			0.81	3000
macro avg	0.73	0.57	0.58	3000
weighted avg	0.78	0.81	0.76	3000

---

[Colab paid products](#) - [Cancel contracts here](#)

0s completed at 21:14

